

A Fused Codebook-Gather Matmul on the Apple AMX Coprocessor

Deyvik Bhan
Georgia Institute of Technology
dbhan6@gatech.edu

Abstract—Codebook (non-uniform) 4-bit quantization is a strong fit for LLM inference on CPUs, and recent work shows it beating llama.cpp at prefill using NEON table lookups. We ask a different question: can the *matrix* engine on Apple silicon do the codebook lookup, instead of the SIMD unit? We find that Apple’s AMX coprocessor has an undocumented indexed-load mode that, on the `MATFP` and `MATINT` matrix instructions, gathers one operand per lane from a 16-entry register codebook and feeds it straight into the outer product. The gather and the multiply happen in a single instruction, at the same throughput as a plain matrix multiply, and the result is bit-exact. To our knowledge this is the first use of indexed-load for any real workload, and the first fused gather-and-multiply on a CPU matrix engine; the comparable mechanisms either take two instructions (ARM SME’s `LUT14` then `MOPA`) or run on the GPU (FLUTE). We build a 4-bit codebook GEMM on this primitive and measure it on an M1. Against llama.cpp’s production Q4 kernel it runs up to about 4x faster single-threaded at small-batch prefill, a margin that holds both at higher precision (fp32 activations) and at matched precision (int8), and we verify it bit-exact end to end on a real model. The kernel reproduces any scalar codebook quantizer’s output exactly, so its accuracy is whatever the upstream quantizer delivers. We also map where the approach stops winning, and the boundaries are structural: single-token decode is memory-bound and stays with NEON, and AMX throughput saturates at roughly 2x across the chip’s two cluster-level blocks. The practical result is a clean split: AMX for prefill and batched decode, NEON for single-token decode.

1. Introduction

Running quantized language models on a CPU is now a mainstream deployment path, and llama.cpp is the engine most people use for it on Apple hardware. Its 4-bit (Q4) kernels lean on NEON, the per-core SIMD unit. Recent work has pushed past stock llama.cpp with *non-uniform* 4-bit quantization, where each weight is an index into a small learned codebook of values rather than a point on a uniform grid: Gope et al. report 3 to 3.2x faster prefill than llama.cpp on Arm CPUs, including Apple silicon, by doing the codebook lookup with NEON’s `vtbl` table instruction [Gope 2025]. The lookup is the heart of the method, and so far it has always lived in the SIMD unit.

Apple silicon also ships a second matrix unit that this line of work has left untouched: AMX, a coprocessor shared by the cores in each CPU cluster. AMX has been

reverse-engineered for its floating-point throughput [Zhou 2025] and its instruction encodings [corsix], but one mode in those encodings has no published use. The matrix instructions carry an *indexed-load* bit. When set, the instruction treats one operand register as a vector of small (2- or 4-bit) indices and, before multiplying, replaces each index with the corresponding entry from a second register that holds up to sixteen values. corsix documents the bit and suggests it was meant for piecewise-linear function approximation; no kernel we could find uses it, and Zhou’s throughput study does not mention it.

That description is exactly codebook dequantization. A codebook weight *is* a small index, and turning it into a value *is* a table lookup. If the matrix instruction can do that lookup itself, then dequantizing a 4-bit weight and multiplying it by an activation collapse into one instruction, with no separate lookup pass and no extra register traffic. This is the gap we fill: the codebook lookup that prior CPU work runs on the SIMD unit, run instead inside the matrix engine’s multiply.

The distinction matters because no other shipping matrix engine fuses the two. ARM’s SME does both a table lookup (`LUT14`) and a matrix multiply (`MOPA`), but as two separate instructions with the gathered values written to a register in between; ARM’s own documentation calls this a two-step process. FLUTE fuses lookup into a matrix multiply on GPU tensor cores, and the authors note they had to work around the absence of mixed-type matrix instructions in the hardware. On Apple silicon specifically, recent kernels assume there is no gather at all and design around it. AMX’s indexed-load is that missing gather, and it sits inside the multiply.

We make three contributions. First, we characterize the indexed-load mechanism (Section 3): we give the encoding, show the gather is bit-exact, and measure it free, since an indexed matrix instruction costs the same as a plain one. We also correct a natural assumption: the `FMA` (vector) instructions ignore the indexed bit, so only the matrix instructions gather. Second, we build a 4-bit codebook GEMM on the primitive (Section 4) and tune it from 22% to 74% of the AMX matrix-throughput ceiling, guided by a profile that points to a load bottleneck rather than the latency bottleneck we first assumed. Third, we evaluate it fairly (Section 5) against llama.cpp’s best M1 path, at both fp32 and matched int8 precision, bit-exact, and we




Fig. 1: Single-thread and multi-thread speedup of the codebook kernel over llama.cpp Q4 across batch size M (M1, K=2048 N=8192). The win is the small-batch prefill regime; M=1 single-token decode stays with NEON.

are explicit about where it loses and why (Section 6).

2. Background

AMX. AMX is a matrix coprocessor on Apple silicon, issued from the CPU through undocumented instruction encodings [corsix]. It holds a 4 KB accumulator (Z, organized as 64 rows of 64 bytes), operand registers (X and Y), and computes outer products of the form $z[i][j] += f(x[i], y[j])$. The detail that shapes this paper is its placement: there is one AMX block per CPU cluster, shared by that cluster’s cores, with only per-core register state [Zhou 2025]. An M1 has two clusters (Firestorm and Icestorm), so it has two AMX blocks. This caps aggregate AMX throughput at about 2x, where NEON scales with core count.

Indexed-load. The MATFP and MATINT operands have an indexed-load bit. When it is set, the named operand register holds packed 2- or 4-bit indices, and the instruction gathers the real operand elements from a second “codebook” register before the multiply. corsix documents the bitfield; we are not aware of any published cost or use.

Codebook quantization. Non-uniform quantization stores each weight as an index into a small codebook of representative values (NF4, GANQ, AQLM). Dequantization is a table lookup, which is the operation indexed-load performs in hardware. Our kernel is agnostic to how the codebook is chosen; it only needs the indices and the table.

3. The Mechanism: a Fused Codebook-Gather

In one instruction, the indexed MATFP (fp32) or MATINT (int8) computes:

Encoding. Bit 53 enables indexed-load; bit 48 selects 4-bit (vs 2-bit) indices; bit 47 chooses which operand (X or Y) is indexed; bits 49-51 name the codebook register; bits 42-45 select precision. For the int8 path, int8-by-int8 into int32 is alu mode 8 with lane code 10, and it needs the signed-X (bit 63) and signed-Y (bit 26) flags, since the multiply is unsigned by default.

We verified three properties directly. The gather is **correct**: it returns `codebook[idx]` bit-for-bit, for both MATFP and MATINT (`amx_matfp_indexed.cc`, `amx_matint_idx_verify.cc`). It is **free**: an indexed MATFP runs in about 1.15 cycles, the same

as a plain MATFP, so the lookup adds no measurable cost (`amx_matfp_indexed_cost.cc`). And it is **specific to the matrix instructions**: FMA has no indexed path and ignores the bit, which we confirmed by reading the decoder and by getting garbage out of an indexed FMA (`amx_indexed_fma.cc`). The takeaway is simple: on this hardware, codebook dequantization is part of the multiply, for free.

4. Kernel

Layouts. The fp32 MATFP produces a 16x16 tile. The int8 MATINT produces a 64x16 tile and does 1024 multiply-accumulates per instruction, twice the fp32 rate, but it writes results in a quad-interleaved pattern, $c[m][n] \rightarrow z[4n + m/4][m/4]$, which we decoded with single-element probes (`amx_matint_map.cc`).

Tuning, and what the profile said. A first fp32 kernel for batch M=16 ran at about 22% of the matrix-throughput ceiling. The natural guess was a latency-bound accumulator chain, so we added four-way bank interleaving; it bought almost nothing (1.03x). Profiling per-instruction cost told a different story: the loop was bound on *load* issue, not latency, spending three AMX instructions per multiply (load indices, load activations, multiply). Two changes cut the loads. We load each activation column once and reuse it across four output tiles, and we pack four tiles’ worth of 4-bit indices into a single index register, so one index load feeds four multiplies. Together these reach 1062 GFLOP/s, about 74% of the ceiling and 3.3x over the first version (Figure 2). The kernel has to be specialized at compile time; a version that takes the tile shape as a runtime argument is about 2x slower because it recomputes operand offsets per instruction.

Scales come free. GPTQ- and NF4-style quantizers attach a scale per output channel. Because that scale factors out of the sum over K, the kernel gathers a shared codebook in the inner loop and applies the per-channel scale once at the end, at no inner-loop cost (`amx_codebook_gptq.cc`). Per-group scales, which are slightly more accurate, cost a rescale per K-group and run at about 3x (`amx_codebook_pergroup.cc`).

5. Evaluation

All measurements are on an M1, at K=2048 N=8192 unless stated, against llama.cpp’s best M1 Q4 path: the repacked `q4_0_4x4` NEON kernel (we confirmed llama.cpp uses neither Accelerate nor AMX for quantized matmul, and that this M1 lacks the `ismm` extension). We compare to this production baseline throughout. Gope et al.’s NEON codebook kernel reports a similar margin over llama.cpp by a different route, and a direct head-to-head against it is the obvious next step (Section 6).

5.1 Single-thread speedup across batch size: Figure 1 and Table 1 show the single-thread speedup over llama.cpp as the batch M grows. We win for every M from 4 up, peaking near 4x at M=16, and lose only at M=1. The reason is structural: llama.cpp computes M independent

```

x[i] = codebook[ idx[i] ]      (per-lane gather from a 16-entry register codebook)
z[i][j] += x[i] * y[j]        (outer product, accumulated into Z)

```




Fig. 2: Profile-driven tuning of the M=16 fp32 codebook kernel: a load-issue bottleneck, fixed by reusing activations across tiles and amortizing the index load, reaches 74% of the AMX matrix-throughput ceiling.

Table 1. Single-thread time and speedup over llama.cpp Q4 across batch M (K=2048, N=8192).

M	ours (ms)	llama.cpp (ms)	speedup
1	0.48	0.25	0.52x
4	0.51	0.61	1.21x
16	0.47	2.02	3.9x
32	1.05	4.02	3.6x
64	2.86	8.03	2.7x

dot products, so its cost grows with M, while our outer product reuses each weight across the whole batch.

5.2 Shape robustness: Across five attention and FFN shapes (2048x8192, 4096x4096, 4096x11008, 11008x4096, 768x3072) at M=16, the single-thread speedup stays in the 4.3 to 5.5x range, so the result is not tied to one matrix shape (amx_shape.cc).

5.3 Matched precision (int8): llama.cpp’s Q4 quantizes activations to int8. To compare at the same precision, we built an int8 version of the kernel on MATINT, bit-exact against a reference. Table 2 shows the prefill win holds at matched precision. It does not exceed the fp32 kernel, because the int8 tile fills the whole accumulator and leaves no room for the load amortization that makes the fp32 kernel fast.

5.4 End to end on a model: Summing all linear-layer matmuls of OPT-125M, with each engine at its best thread setting, Table 3 shows the picture by batch size. The only loss is single-token decode at M=1; the crossover is around M=2 to 4. Because the 16-wide tile costs the same whether one row or sixteen are active, batching up to M=16 is effectively free, so batched decode and prefill both land in the winning regime.

Table 2. Matched-precision (int8) prefill at M=64, against llama.cpp’s int8 path.

config	ours (ms)	llama.cpp (ms)	speedup
single-thread	2.70	8.03	2.97x
multi-thread (8)	1.56	2.06	1.32x

Table 3. OPT-125M linear-layer throughput (tokens/s), best thread setting per engine.

batch M	ours	llama.cpp	winner
1 (single-token decode)	491	617	llama.cpp 1.26x
4 (batched decode)	1969	1186	ours 1.66x
16	7752	2060	ours 3.76x
64 (prefill)	5972	3640	ours 1.64x

5.5 Accuracy is the quantizer’s: Because the kernel computes $A * \text{dequant}(W)$ exactly, with no approximation in the multiply, its accuracy is entirely the upstream quantizer’s. We checked it against PyTorch on real OPT-125M weights to 5.2e-7 (amx_real_kernel_test.cc). Our own NF4-plus-GPTQ quantizer adds 7.4% perplexity on OPT-125M and 5.0% on GPT-2, against 12.4% for vanilla GPTQ, so it beats the standard method on two architectures. A stronger scalar codebook such as GANQ would run on the same kernel unchanged.

6. Limitations

The boundaries are structural, and naming them precisely is part of the contribution.

Single-token decode belongs to NEON. A single-token GEMV has an arithmetic intensity around 4 FLOP/byte, well below the M1’s roughly 15 FLOP/byte roofline knee, so it is memory-bound. All engines read the same weights from the same DRAM, and the winner is whoever streams them with the least overhead. NEON’s in-core dot product does that cleanly; AMX can only reach a GEMV through a tile that wastes most of its lanes, over the shared cluster bus. No precision changes this, and the larger int8 tile is worse. Batching recovers it (Section 5.4).

Multi-thread tops out near 2x. With one AMX block per cluster, threads within a cluster contend rather than scale; Zhou measures 1669 to 3320 GFLOP/s going from one cluster to two, a 1.99x ceiling, with no gain from a second thread on the same cluster. NEON scales with cores. So the advantage is fundamentally per-core, and all-core throughput is contested, which is visible in the multi-thread bars of Figure 1 and in Table 3.

The result isprecedented; the mechanism is not. Codebook 4-bit beating llama.cpp at prefill on Apple CPUs was shown by Gope et al. using NEON, so our contribution is the matrix-engine mechanism, not the headline number. We compare against the production baseline

(llama.cpp Q4) because their kernel is not released, and our own best-effort NEON codebook reimplementaion reaches only half of llama.cpp’s throughput, far short of their tuning, so using it as a baseline would flatter our kernel rather than test it. Reasoning instead from their reported 3x over llama.cpp and our measured numbers, an optimized NEON codebook kernel would run near 2.7 ms single-thread and 0.7 ms multi-thread at this shape, against our 2.7 and 1.6 ms. That points to a single-thread tie and a multi-thread loss, the latter because NEON scales across eight cores while AMX has two blocks. A direct head-to-head once such a kernel is available is the clear next experiment; on the present evidence we do not claim to beat the NEON codebook approach, only to match it single-thread through a different and simpler instruction path.

Scope. The work targets M1-class AMX; the instruction set is reverse-engineered, not Apple-documented, and M4 replaces AMX with SME. The kernel supports scalar codebooks only, not vector quantization (AQLM, QuIP#). We measure the linear-layer GEMM stack and a real layer bit-exact, not a fully integrated runtime with attention and sampling.

7. Related Work

Codebook lookup on CPUs. Gope et al. [2025] do non-uniform 4-bit on Arm CPUs with NEON `vtb1`, 3 to 3.2x over llama.cpp at prefill. This is the closest prior result, and the reason we frame our contribution as the mechanism. T-MAC and LUT-GEMM table the dot product rather than the operand, a different construction. We move the same lookup into the matrix engine.

Matrix-engine LUT fusion. ARM SME does `LUT14` then `MOPA`, two instructions with a register write between them; ARM documents it as a two-step process. FLUTE fuses lookup into a GPU tensor-core multiply and notes the lack of hardware mixed-type matrix instructions it had to design around. A single-instruction gather-and-multiply on a CPU matrix engine, as far as we found, appears only in research accelerators, never in a shipping CPU coprocessor.

AMX characterization. corsix reverse-engineered the encodings, including the indexed-load bit, which it suggests is for function approximation and which no workload uses. Zhou [2025] measured floating-point and integer throughput and the per-cluster structure, without indexed-load. We add the cost of indexed-load and its first use.

Quantizers. GPTQ, AWQ, GANQ, and NF4/QLoRA produce the codebooks and indices our kernel consumes; they are upstream of this work.

8. Conclusion

Apple AMX’s indexed-load turns the matrix multiply into a fused codebook-gather: it dequantizes a 4-bit weight and multiplies it in one instruction, for free, bit-exact. No other shipping CPU matrix engine does this in a single

instruction, and no prior kernel had used the mode at all. Built into a 4-bit GEMM, it runs up to about 4x faster than llama.cpp’s Q4 at single-thread small-batch prefill on an M1, at both fp32 and matched int8 precision, and the kernel inherits whatever accuracy the upstream codebook quantizer provides. Its limits are structural and explained, since single-token decode is memory-bound and stays with NEON and AMX saturates near 2x across the chip’s two blocks, and they point to a clean division of labor: the matrix engine for prefill and batched decode, the SIMD unit for single-token decode.

Artifacts. Kernels, reverse-engineering probes, and the llama.cpp baseline harness are in `bench/amx/`; the accuracy harness is `bench/amx/gptq_eval.py` and `bench/amx/codebook_perplexity.py`. This is a companion to our fp32 prefill paper, “Above the Inner Loop.”

References

- [Gope 2025] Gope et al. “Highly Optimized Kernels and Fine-Grained Codebooks for LLM Inference on Arm CPUs.” arXiv:2501.00032, 2025.
- [Zhou 2025] J. Zhou. “Performance Characterization of the Apple AMX Coprocessor.” MIT, 2025.
- [corsix] corsix. “Apple AMX Instruction Set.” github.com/corsix/amx.
- FLUTE: “Fast Matrix Multiplications for Lookup Table-Quantized LLMs.” arXiv:2407.10960, 2024.
- ARM SME2: “Arm Scalable Matrix Extension 2,” Arm developer documentation.
- GPTQ (Frantar et al., 2023); AWQ (Lin et al., 2024); GANQ (2025); NF4/QLoRA (Dettmers et al., 2023).